

# Lecture 1: Introduction to Spark and HPC

---

[Haiping Lu](#)

[COM6012: Scalable ML](#)

YouTube Playlist:

<https://www.youtube.com/c/HaipingLu/>



# Week 1 Contents / Objectives

- The Big Data Problem: Why Spark?
- What is Spark?: The Essentials
- An Example of Spark: Log Mining
- How to Use Spark: PySpark, HPC, Resources

# Week 1 Contents / Objectives

- The Big Data Problem: Why Spark?
- What is Spark?: The Essentials
- An Example of Spark: Log Mining
- How to Use Spark: PySpark, HPC, Resources

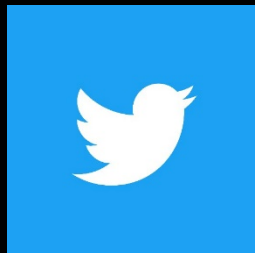
# Where Does Big Data Come From?

- All happening **online**, e.g. tracking of:
  - Clicks
  - Billing events
  - Server requests
  - Transactions
  - Network messages
  - Faults
  - ...

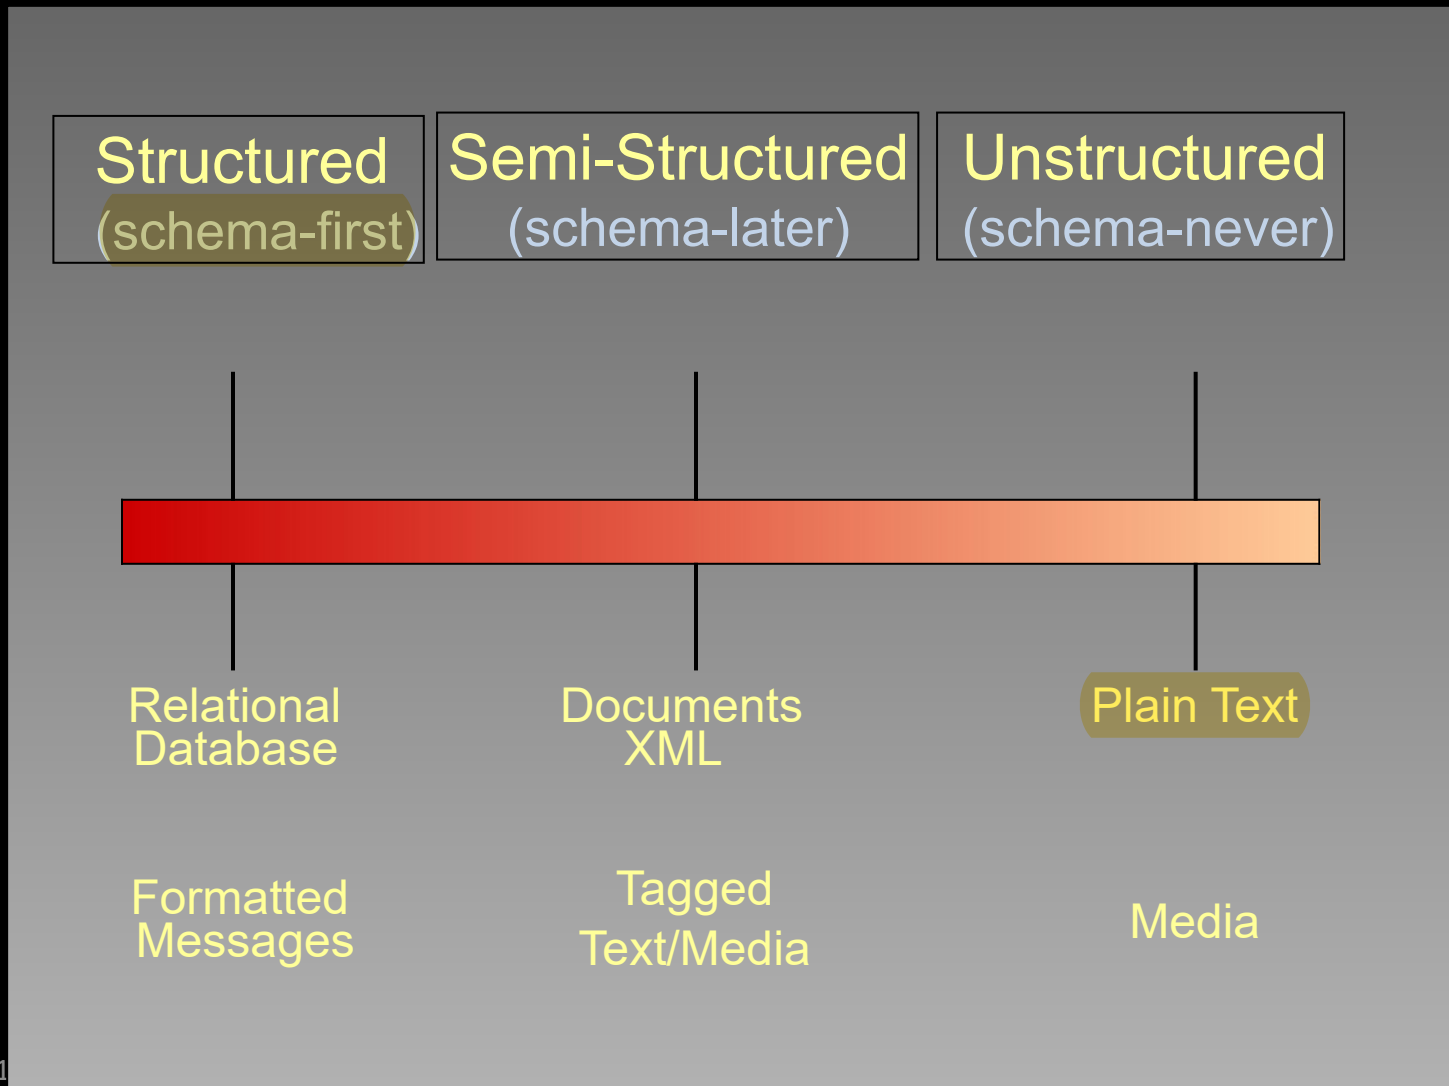


# Where Does Big Data Come From?

- User generated content: web + mobile

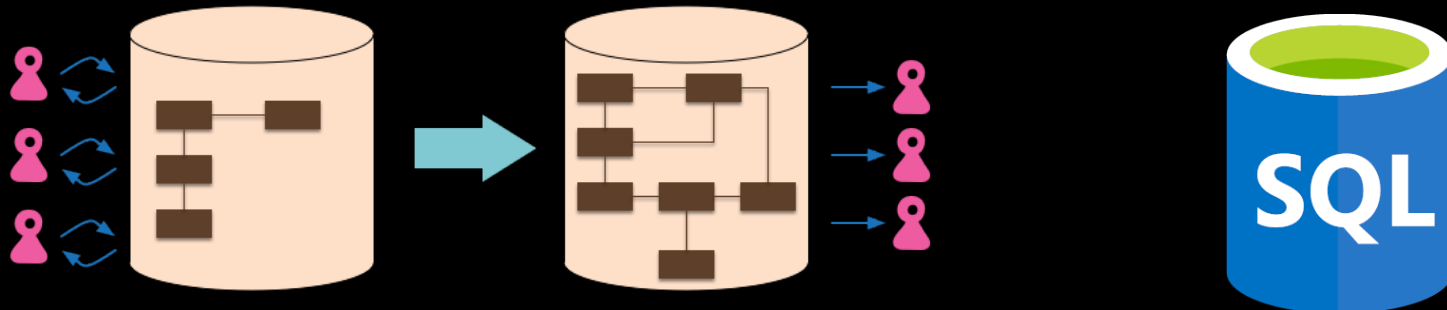


# Data Structure Spectrum



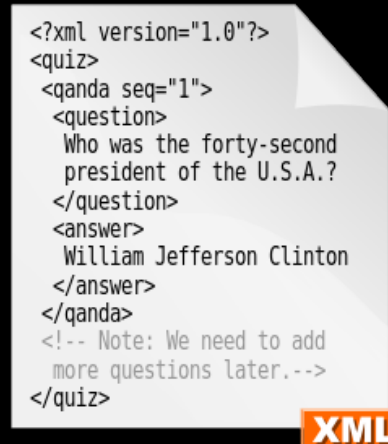
# Structured Data

- **Database:** relational data model → how a database is structured and used (hottest job 20 years ago)
- **Schema:** the organisation of data as a blueprint of how the database is constructed
  - The programmer **must statically specify** the schema
  - Decreasing ← consumer/media app, enterprise search
- **SQL:** Structured Query Language



# Semi-Structured Data

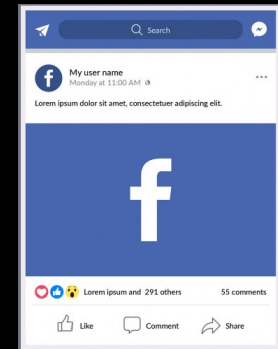
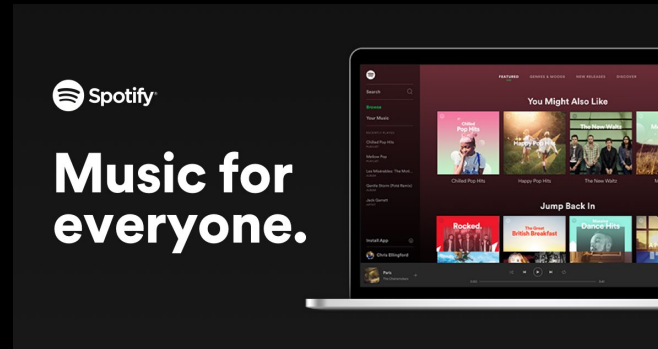
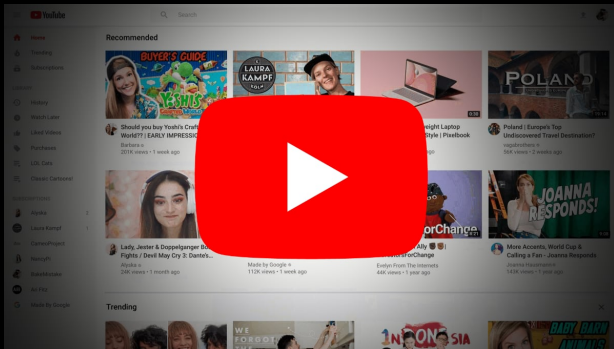
- **Self-describing** rather than formal structures, tags/markers to separate semantic elements
- The column types → the **schema** for the data
  - **Spark** dynamically infers the schema while reading each row
  - Programmer statically specifies the schema
- Examples:





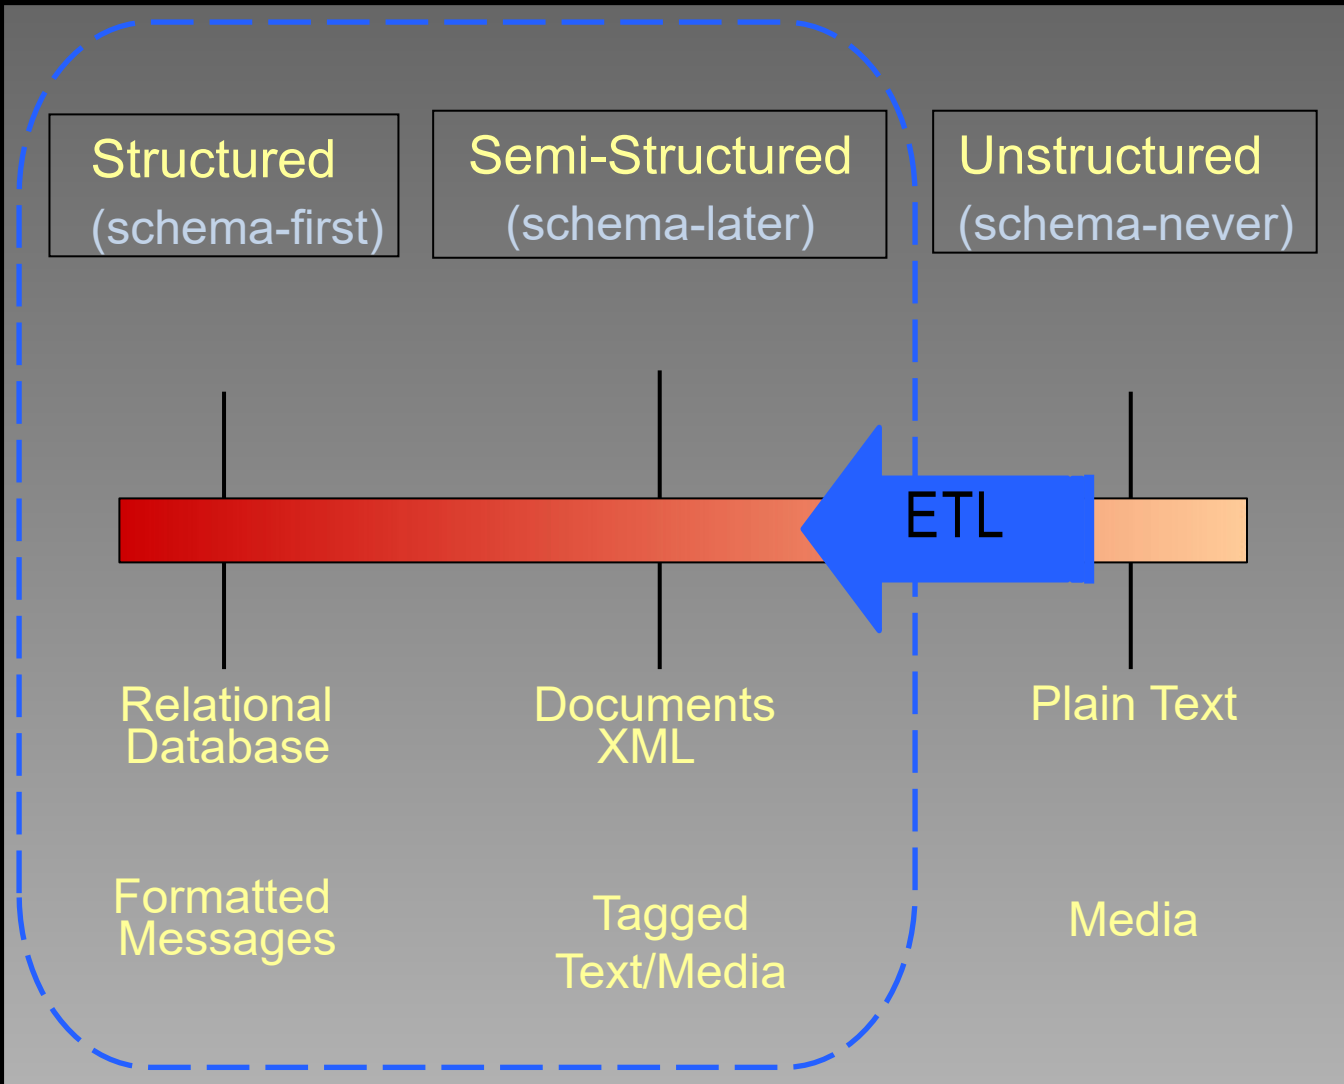
# Unstructured Data

- Only one column with string or binary type
- Examples



- More than 70%–80% of all data in organisations (Shilakes 1998)

# Traverse the Data Structure Spectrum



- Impose structure on unstructured data

- Extract
- Transform
- Load

# Traditional Analysis Tools

- Unix shell commands (awk, grep, ...)

```
root@enginx:~# awk ' {print $0}' file.txt
```


Item	Model	Country	Cost
1	BMW	Germany	\$25000
2	Volvo	Sweden	\$15000
3	Subaru	Japan	\$2500
4	Ferrari	Italy	\$2000000
5	SAAB	USA	\$3000

```
vulphere@arifuretaarch:~|⇒ grep root /etc/passwd
root:x:0:0:root:/root:/bin/zsh
vulphere@arifuretaarch:~|⇒ grep -n root /etc/passwd
1:root:x:0:0:root:/root:/bin/zsh
vulphere@arifuretaarch:~|⇒ grep -c false /etc/passwd
3
vulphere@arifuretaarch:~|⇒ _
```

All run on a single machine!


# The Big Data Problem

- Data growing faster than computation speeds
- Growing data sources
  - Web, mobile, scientific, ...
- Storage getting cheaper
  - Size doubling every 18 months
- But, stalling CPU speeds and storage bottlenecks




# Solution for the Big Data Problem

- **One machine** can not process or *even store* all the data!
- Solution: **distribute** data over a **cluster** of machines



Lots of hard drives



... and CPUs



... and memory!

# Week 1 Contents / Objectives

- The Big Data Problem: Why Spark?
- What is Spark?: The Essentials
- An Example of Spark: Log Mining
- How to Use Spark: PySpark, HPC, Resources

# Apache Spark

- Fast and general **cluster** computing system

- Interoperable with 

- Improves efficiency through:

- **In-memory** computing primitives
- General **computation** graphs

→ Up to 100× faster  
(2-10× on disk)

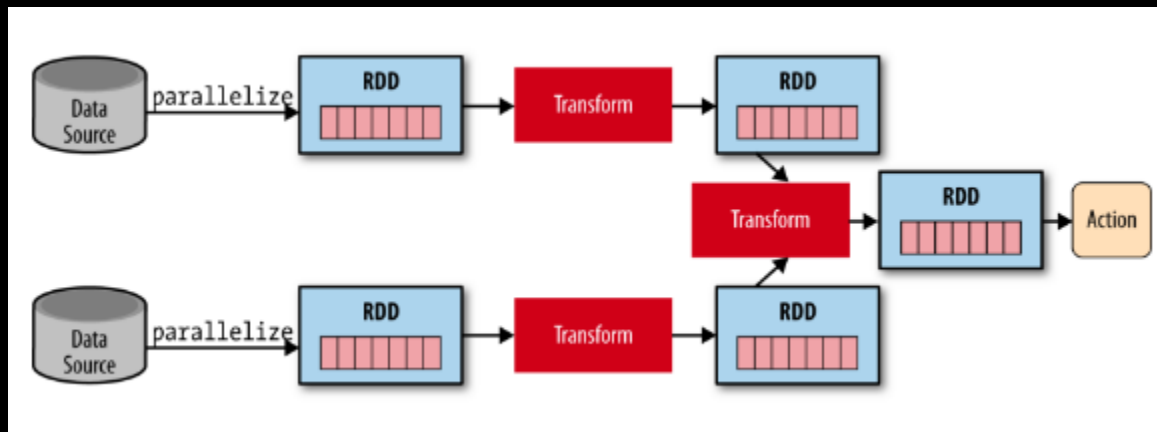
- Improves usability through:

- Rich APIs in Scala, Java, **Python**
- **Interactive shell**

→ 2-5× less code

# Spark Model

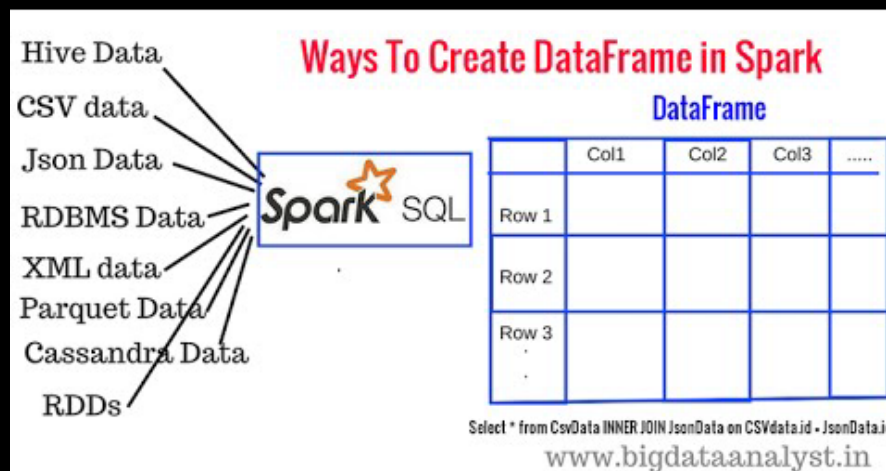
- Write programs in terms of **transformations** on **distributed** datasets
- Resilient Distributed Datasets (RDDs)
  - **Collections** of objects that can be stored in memory or disk across a cluster
  - **Parallel** functional transformations (map, filter, ...)
  - Automatically rebuilt on **failure**





# Spark for Data Science

- DataFrames
  - **Structured** data (**SQL**)
  - Familiar API based on R/Python Pandas
  - Distributed, optimised implementation
- Machine learning pipelines
  - Simple construction and tuning of **ML workflows**



# Spark Computing Framework

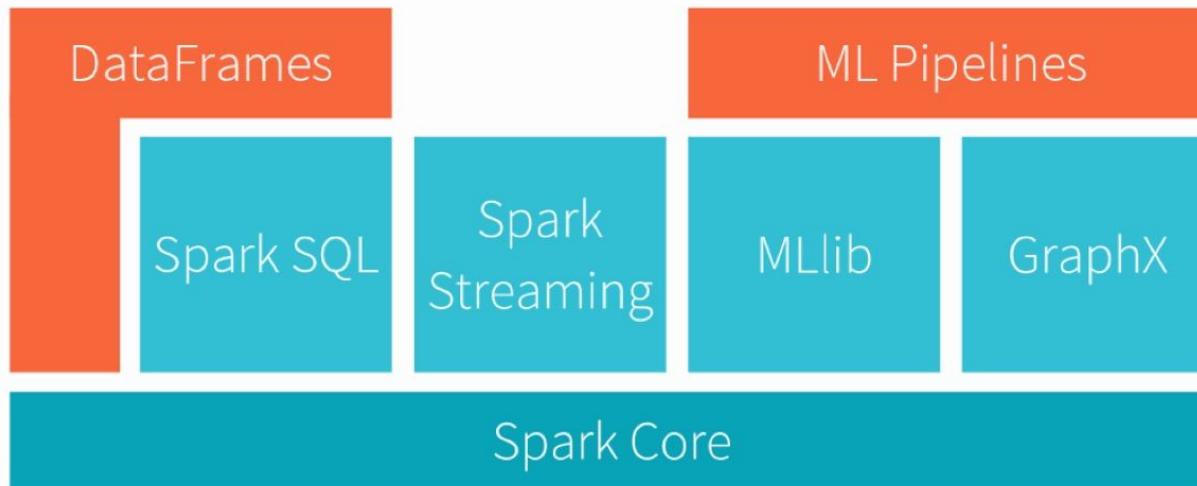
- Programming abstraction and parallel runtime to hide complexities of **fault-tolerance** and **slow** machines

“Here’s an operation, run it on all of the data”

**JUST DO IT.**

- I don’t care where it runs (you schedule that)
- In fact, feel free to run it twice on different nodes (e.g. when it fails)

# Apache Spark Ecosystem




Total contributors: 150 → 500

Lines of code: 190K → 370K

500+ active production deployments

<https://i.pinimg.com/originals/e7/f3/2d/e7f32d041846a5938a09e192bdf3885d.jpg>

# Spark Components




- A Spark program first creates a **SparkSession** object as the driver (including **SparkContext**)
  - Tells Spark how/where to access a cluster
  - Connect to cluster managers
- Cluster managers
  - Allocate resources across applications
- Spark executor (worker):
  - Run computations
  - Access data storage

# SparkSession and SparkContext

- SparkSession
  - Entry point for DataFrame API, create **DataFrames**
  - PySpark shell automatically create SparkSession as **spark**
  - Programs: must create a new SparkSession first (see lab)
- SparkContext
  - Entry point for Spark functionality, create **RDDs**
  - Connect to a Spark cluster
  - Associated with a SparkSession
  - PySpark shell automatically create SparkContext as **sc**
  - Programs: **sc = spark.sparkContext**

# The 'Master' Parameter for a SparkSession

- Determines cluster type and size

Master Parameter	Description
<code>local</code>	run Spark locally with one worker thread (no parallelism)
<code>local[K]</code> 	run Spark locally with K worker threads (ideally set to number of cores)
<code>spark://HOST:PORT</code>	connect to a Spark standalone cluster; PORT depends on config (7077 by default)
<code>mesos://HOST:PORT</code>	connect to a Mesos cluster; PORT depends on config (5050 by default)

# Week 1 Contents / Objectives

- The Big Data Problem: Why Spark?
- What is Spark?: The Essentials
- An Example of Spark: Log Mining
- How to Use Spark: PySpark, HPC, Resources


# Spark Example: Log Mining (w/t RDD)

Load error messages from a log into memory, then interactively search for various patterns



# Spark Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns



# Spark Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")
```



# Spark Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

Base RDD

```
lines = spark.textFile("hdfs://...")
```

Driver

Worker

Worker

Worker

# Spark Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")  
errors = lines.filter(lambda s: s.startswith("ERROR"))
```



# Spark Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

Transformed RDD

```
lines = spark.textFile("hdfs://...")  
errors = lines.filter(lambda s: s.startswith("ERROR"))
```

Driver

Worker

Worker


Worker

# Spark Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")
errors = lines.filter(lambda s: s.startswith("ERROR"))
messages = errors.map(lambda s: s.split("\t")[2])
messages.cache()
```

```
messages.filter(lambda s: "mysql" in s).count()
```



# Spark Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")
errors = lines.filter(lambda s: s.startswith("ERROR"))
messages = errors.map(lambda s: s.split("\t")[2])
messages.cache()
```

```
messages.filter(lambda s: "mysql" in s).count()
```

Driver

Action

Worker

Worker

Worker

# Spark Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")
errors = lines.filter(lambda s: s.startswith("ERROR"))
messages = errors.map(lambda s: s.split("\t")[2])
messages.cache()
```

```
messages.filter(lambda s: "mysql" in s).count()
```

Driver

Worker

Block 1

Worker

Block 3

Worker

Block 2

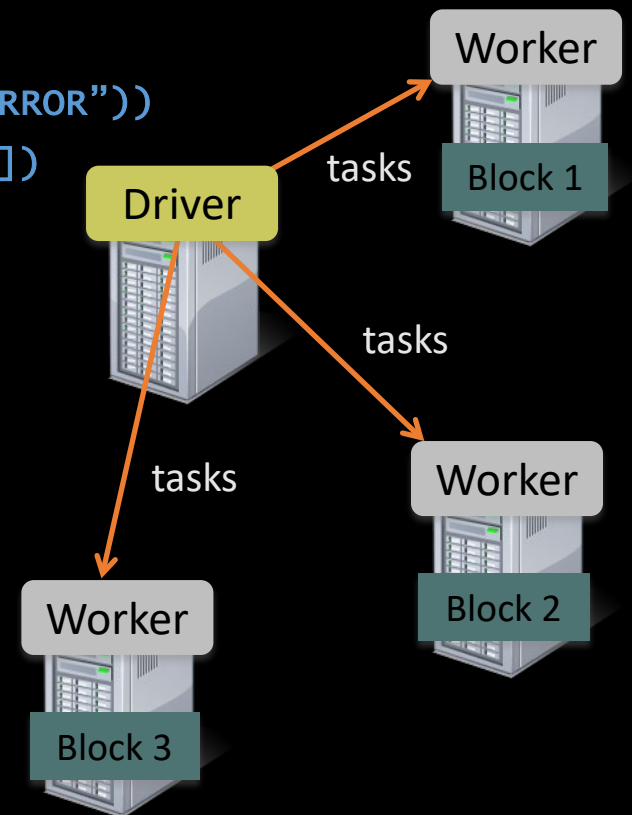


# Spark Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")
errors = lines.filter(lambda s: s.startswith("ERROR"))
messages = errors.map(lambda s: s.split("\t")[2])
messages.cache()
```

```
messages.filter(lambda s: "mysql" in s).count()
```

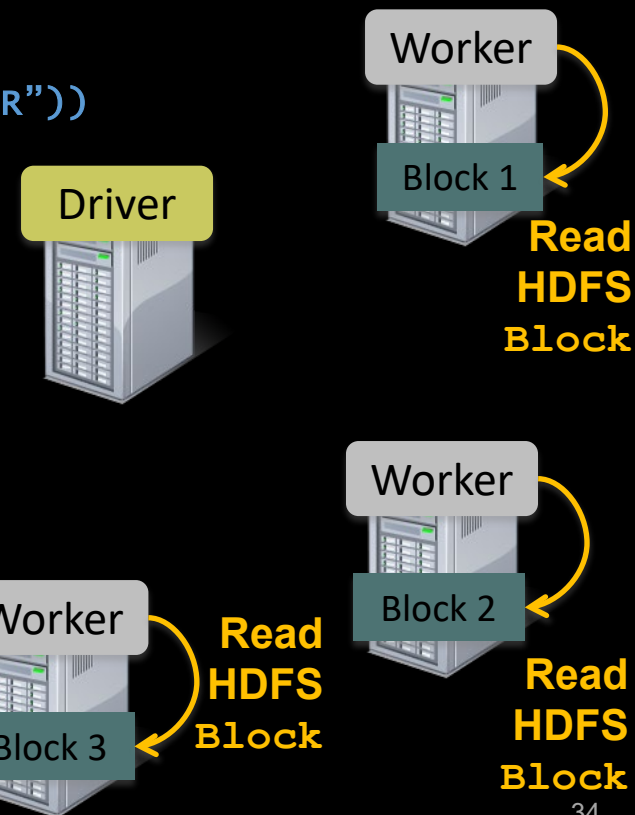


# Spark Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")
errors = lines.filter(lambda s: s.startswith("ERROR"))
messages = errors.map(lambda s: s.split("\t")[2])
messages.cache()
```

```
messages.filter(lambda s: "mysql" in s).count()
```

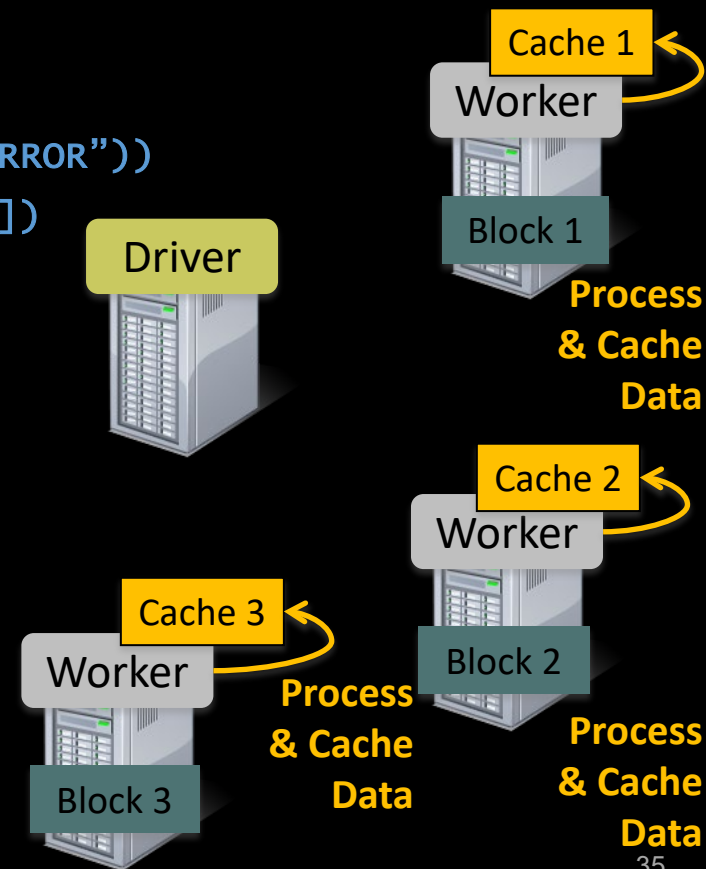


# Spark Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")
errors = lines.filter(lambda s: s.startswith("ERROR"))
messages = errors.map(lambda s: s.split("\t")[2])
messages.cache()
```

```
messages.filter(lambda s: "mysql" in s).count()
```




# Spark Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")
errors = lines.filter(lambda s: s.startswith("ERROR"))
messages = errors.map(lambda s: s.split("\t")[2])
messages.cache()
```

```
messages.filter(lambda s: "mysql" in s).count()
```




# Spark Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")
errors = lines.filter(lambda s: s.startswith("ERROR"))
messages = errors.map(lambda s: s.split("\t")[2])
messages.cache()
```

```
messages.filter(lambda s: "mysql" in s).count()
messages.filter(lambda s: "php" in s).count()
```

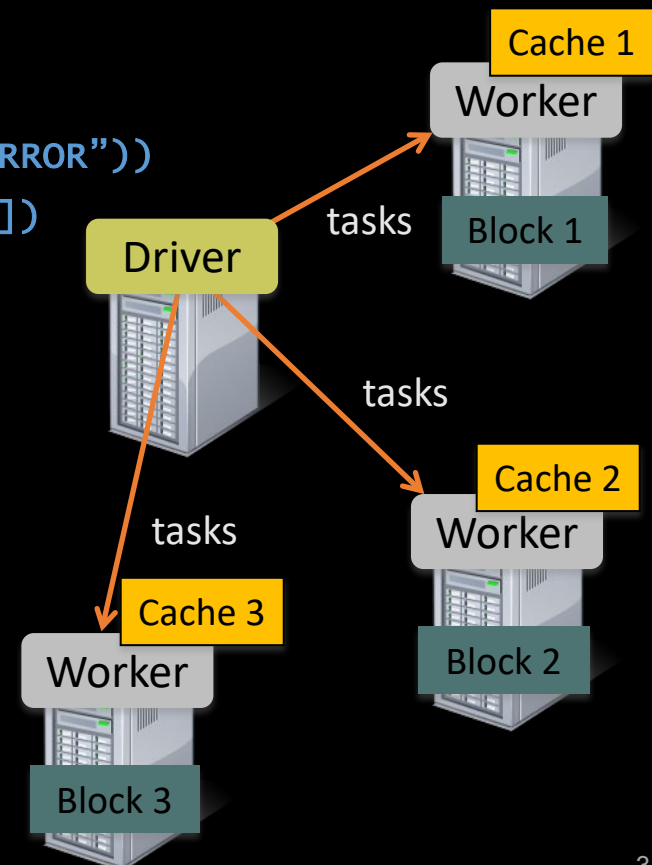


# Spark Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")
errors = lines.filter(lambda s: s.startswith("ERROR"))
messages = errors.map(lambda s: s.split("\t")[2])
messages.cache()
```

```
messages.filter(lambda s: "mysql" in s).count()
messages.filter(lambda s: "php" in s).count()
```

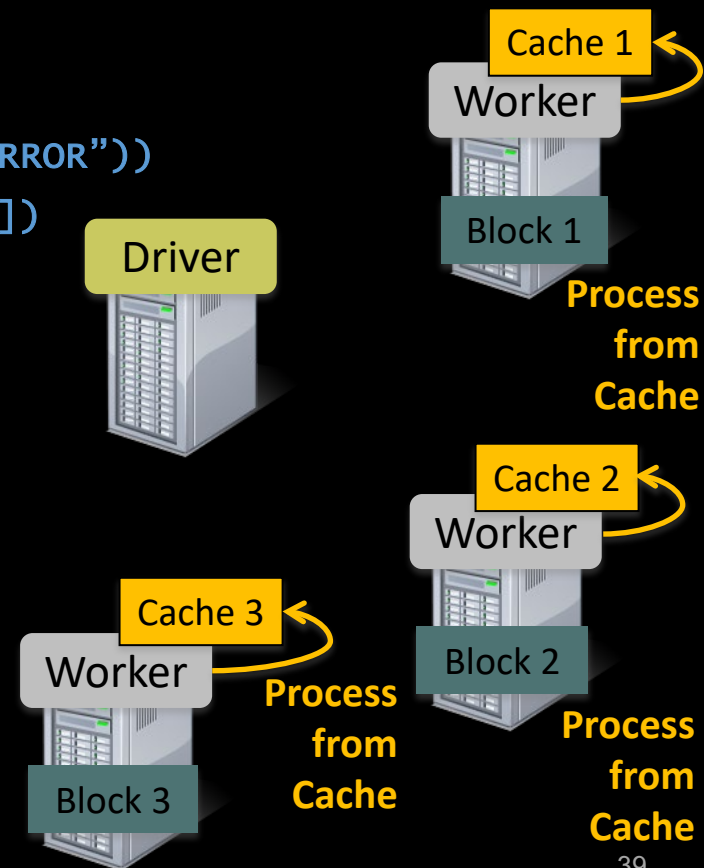


# Spark Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")
errors = lines.filter(lambda s: s.startswith("ERROR"))
messages = errors.map(lambda s: s.split("\t")[2])
messages.cache()
```

```
messages.filter(lambda s: "mysql" in s).count()
messages.filter(lambda s: "php" in s).count()
```




# Spark Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")
errors = lines.filter(lambda s: s.startswith("ERROR"))
messages = errors.map(lambda s: s.split("\t")[2])
messages.cache()
```

```
messages.filter(lambda s: "mysql" in s).count()
messages.filter(lambda s: "php" in s).count()
```





# Spark Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns


```
lines = spark.textFile("hdfs://...")
errors = lines.filter(lambda s: s.startswith("ERROR"))
messages = errors.map(lambda s: s.split("\t")[2])
messages.cache()
```

```
messages.filter(lambda s: "mysql" in s).count()
messages.filter(lambda s: "php" in s).count()
```

**Cache** your data → Faster results

*Full-text search of Wikipedia*

- 60GB on 20 EC2 machines
- 0.5 sec from mem vs. 20s for on-disk



# Week 1 Contents / Objectives

- The Big Data Problem: Why Spark?
- What is Spark?: The Essentials
- An Example of Spark: Log Mining
- How to Use Spark: PySpark, HPC, Resources

# Spark Program Lifecycle

- Create DataFrames from external data or [createDataFrame](#) from a collection in a driver program
- Lazily transform them into new DataFrames
- `cache( )` some DataFrames for reuse
- Perform actions to execute parallel computation and produce results

Use Spark Transformations and Actions wherever possible: Search [DataFrame reference API](#)


# PySpark 3.0.1

- Need: **Java**, Python, Spark
- See lab 1 on how to install on HPC
- To install on Windows
  - [Lab 1 instructions](#): Install Java JRE, Python, Spark
  - Or pip install pyspark==3.0.1
- To install on Linux/Mac: see lab references

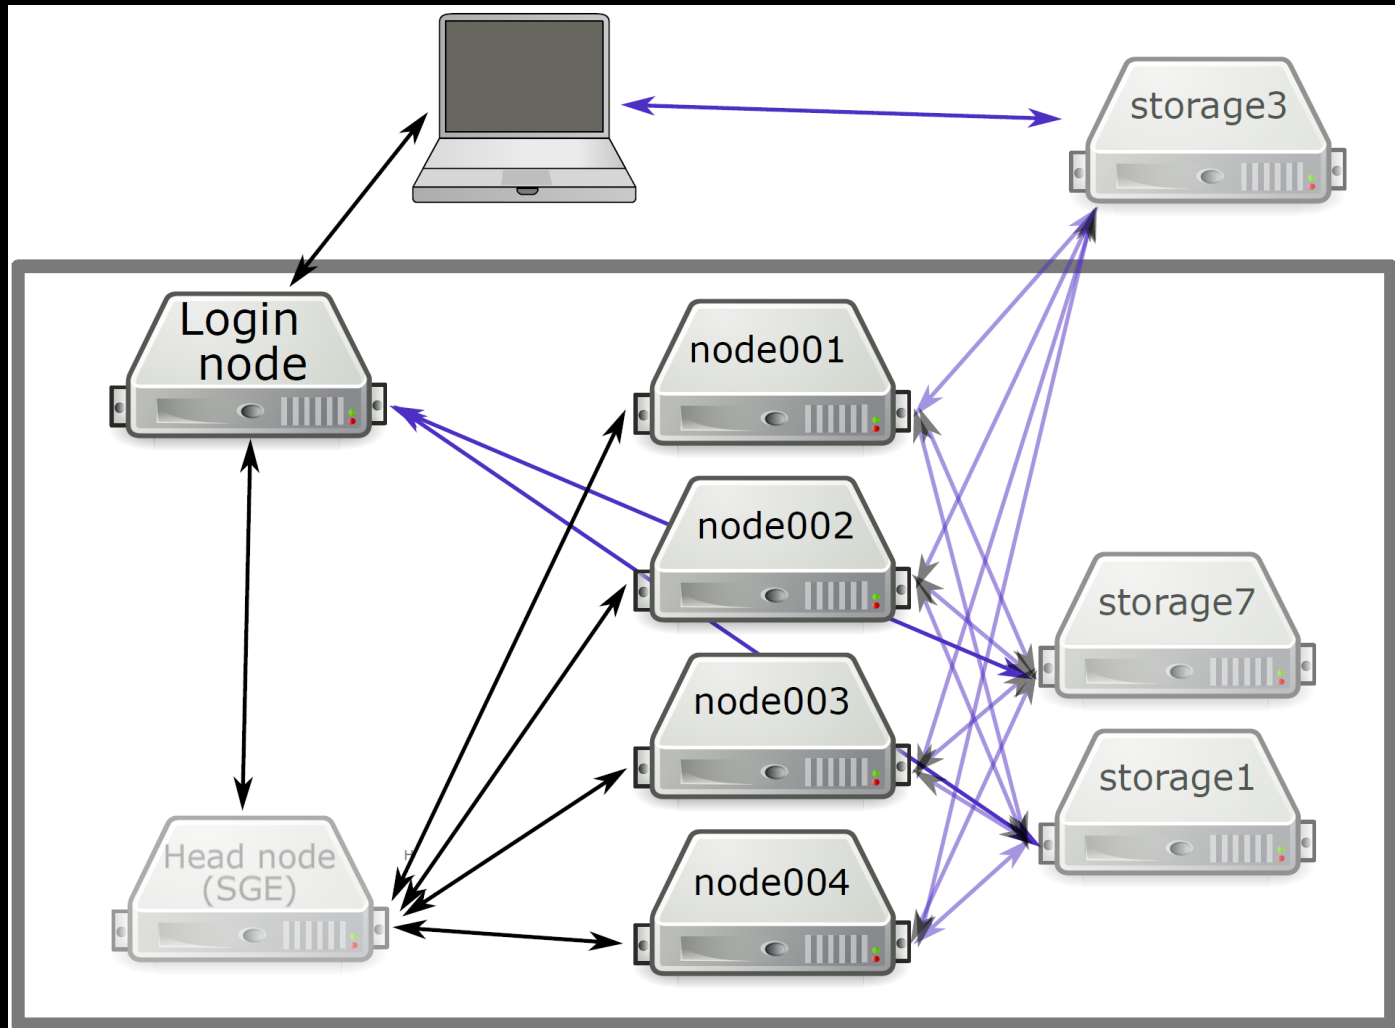


# ShARC HPC @ Sheffield

- ShARC: Sheffield Advanced Research Computer
- VPN: a **MUST** unless you are on a **campus network**
- SSH access via **sharc.sheffield.ac.uk**
  - Windows: MobaXTerm
  - Linux/MAC OS: terminal (command line)
- Help: [hpc@sheffield.ac.uk](mailto:hpc@sheffield.ac.uk)



# HPC Cluster Structure



# Storage

Location	Shared	Quota	Back ups	Speed	Suitable for?
/home/\$USER	Y	10GB	Y	>	Personal data
/data/\$USER	Y	100GB	Y	>	Personal data
/fastdata/\$USER	Y	-	N	>>>	Temporary big files
/scratch	N	-	N	>>>	Temporary small files



# Interactive Session

pyspark

You should see spark version 3.0.1 displayed like below

• • • • •

# Welcome to

```

      _____
     /   _/   _   _   _   _/_/   _/
    \  V/  _  V/  _  `/_/_/  '/_/
   /_/_/ ._/_\_,_/_/_/_/_/\_\
        /_/_/

```

version 3.0.1

Using Python version 3.6.2 (default, Jul 20 2017 13:51:32)

```
SparkSession available as 'spark'.
```

&gt;&gt;&gt;



# Batch Session – Shell Script xx.sh

Create a file `Lab1_SubmitBatch.sh`

```
#!/bin/bash
#$ -l h_rt=6:00:00 #time needed
#$ -pe smp 2 #number of cores
#$ -l rmem=8G #number of memory
#$ -o ../Output/COM6012_Lab1.txt #This is where your output and errors are logged.
#$ -j y # normal and error outputs into a single file (the file above)
#$ -M youremail@shef.ac.uk #Notify you by email, remove this line if you don't like
#$ -m ea #Email you when it finished or aborted
#$ -cwd # Run job from current directory

module load apps/java/jdk1.8.0_102/binary

module load apps/python/conda

source activate myspark

spark-submit ../Code/LogMiningBig.py # .. is a relative path, meaning one level up
```

# Batch Session: Submit & Relax

- **qsub** your job (can run at the login node): see Lab 1
- Then?
  - Close the terminal and leave
  - Wait for pre-set email notification
  - Check status: **qstat**
  - Cancel/amend job: **qdel**
- How much resources to request
  1. Run **short** test jobs
  2. View resource utilisation
  3. Extrapolate
  4. Submit larger jobs



# Spark Resources

- [Apache Spark Documentation](#) (3.0.1)
- [PySpark tutorial](#) (last update Feb 2020)
- [Spark videos on YouTube](#)
- [Open source code](#)
- Suggested reading in labs

## Suggested reading:

- [Spark Overview](#)
- [Spark Quick Start](#) (Choose **Python** rather than the default *Scala*)
- Chapters 2 to 4 of [PySpark tutorial](#) (several sections in Chapter 3 can be safely skipped)
- Reference: [PySpark documentation](#)
- Reference: [PySpark source code](#)

# Acknowledgements

- Some slides (sec. 1) are modified from the “[Introduction to Apache Spark](#)” course by Prof. A. D. Joseph, University of California, Berkeley.
- This module benefits from many open resources. See the acknowledgement on our [GitHub page](#).
- There are many other resources that I have consulted but somehow lost track of the origins.

